

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Lasso
from sklearn import metrics
```

```
In [2]: car_dataset = pd.read_csv('car data.csv')
```

```
In [3]: car_dataset.head()
```

```
Out[3]:
```

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	On
0	ritz	2014	3.35	5.59	27000	Petrol	Dealer	Manual	
1	sx4	2013	4.75	9.54	43000	Diesel	Dealer	Manual	
2	ciaz	2017	7.25	9.85	6900	Petrol	Dealer	Manual	
3	wagon r	2011	2.85	4.15	5200	Petrol	Dealer	Manual	
4	swift	2014	4.60	6.87	42450	Diesel	Dealer	Manual	

```
In [4]: car_dataset.shape
```

```
Out[4]: (301, 9)
```

```
In [5]: car_dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 301 entries, 0 to 300
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Car_Name        301 non-null   object
1   Year            301 non-null   int64
2   Selling_Price   301 non-null   float64
3   Present_Price   301 non-null   float64
4   Kms_Driven      301 non-null   int64
5   Fuel_Type       301 non-null   object
6   Seller_Type     301 non-null   object
7   Transmission    301 non-null   object
8   Owner           301 non-null   int64
dtypes: float64(2), int64(3), object(4)
memory usage: 21.3+ KB
```

```
In [6]: car_dataset.describe()
```

Out[6]:

	Year	Selling_Price	Present_Price	Kms_Driven	Owner
count	301.000000	301.000000	301.000000	301.000000	301.000000
mean	2013.627907	4.661296	7.628472	36947.205980	0.043189
std	2.891554	5.082812	8.644115	38886.883882	0.247915
min	2003.000000	0.100000	0.320000	500.000000	0.000000
25%	2012.000000	0.900000	1.200000	15000.000000	0.000000
50%	2014.000000	3.600000	6.400000	32000.000000	0.000000
75%	2016.000000	6.000000	9.900000	48767.000000	0.000000
max	2018.000000	35.000000	92.600000	500000.000000	3.000000

In [7]: `car_dataset.isnull().sum()`

Out[7]:

```

Car_Name      0
Year          0
Selling_Price 0
Present_Price 0
Kms_Driven    0
Fuel_Type     0
Seller_Type   0
Transmission  0
Owner         0
dtype: int64

```

In [8]: *# Checking the distribution of categorical data*

In [9]:

```

print(car_dataset.Fuel_Type.value_counts())
print(car_dataset.Seller_Type.value_counts())
print(car_dataset.Transmission.value_counts())

```

```

Petrol      239
Diesel      60
CNG         2
Name: Fuel_Type, dtype: int64
Dealer      195
Individual  106
Name: Seller_Type, dtype: int64
Manual      261
Automatic   40
Name: Transmission, dtype: int64

```

Encoding the Categorical Data

In [10]:

```

# encoding 'Fuel_Type' column
car_dataset.replace({'Fuel_Type': {'Petrol': 0, 'Diesel': 1, 'CNG': 2}}, inplace=True)

# encoding 'Seller_Type' column
car_dataset.replace({'Seller_Type': {'Dealer': 0, 'Individual': 1}}, inplace=True)

# encoding 'Transmission' column
car_dataset.replace({'Transmission': {'Manual': 0, 'Automatic': 1}}, inplace=True)

```

In [11]: `car_dataset.dtypes`

```
Out[11]: Car_Name      object
Year          int64
Selling_Price float64
Present_Price float64
Kms_Driven    int64
Fuel_Type     int64
Seller_Type   int64
Transmission  int64
Owner         int64
dtype: object
```

```
In [12]: car_dataset.head()
```

```
Out[12]:
```

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner_Type
0	ritz	2014	3.35	5.59	27000	0	0	0	0
1	sx4	2013	4.75	9.54	43000	1	0	0	0
2	ciaz	2017	7.25	9.85	6900	0	0	0	0
3	wagon r	2011	2.85	4.15	5200	0	0	0	0
4	swift	2014	4.60	6.87	42450	1	0	0	0

Splitting the data into Training data and test data

```
In [13]: X = car_dataset.drop(['Car_Name', 'Selling_Price'], axis=1)
Y = car_dataset['Selling_Price']
```

```
In [14]: print(X)
print(Y)
```

	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	\
0	2014	5.59	27000	0	0	0	
1	2013	9.54	43000	1	0	0	
2	2017	9.85	6900	0	0	0	
3	2011	4.15	5200	0	0	0	
4	2014	6.87	42450	1	0	0	
..	
296	2016	11.60	33988	1	0	0	
297	2015	5.90	60000	0	0	0	
298	2009	11.00	87934	0	0	0	
299	2017	12.50	9000	1	0	0	
300	2016	5.90	5464	0	0	0	

	Owner
0	0
1	0
2	0
3	0
4	0
..	...
296	0
297	0
298	0
299	0
300	0

[301 rows x 7 columns]

0	3.35
1	4.75
2	7.25
3	2.85
4	4.60
..	...
296	9.50
297	4.00
298	3.35
299	11.50
300	5.30

Name: Selling_Price, Length: 301, dtype: float64

Splitting Training and test data

```
In [15]: X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.1,random_state=2)
```

Model Training

1.Linear Regression

```
In [16]: # Loading the linear regression model

lin_reg_model = LinearRegression()
```

```
In [17]: lin_reg_model.fit(X_train,Y_train)
```

```
Out[17]: LinearRegression()
```

Model Evaluation

```
In [18]: # Prediction on Training data
training_data_prediction = lin_reg_model.predict(X_train)
```


```
In [19]: # R squared Error

error_score = metrics.r2_score(Y_train,training_data_prediction)
print("R squared Error : ",error_score)
```

R squared Error : 0.8799451660493705

Visualize the actual prices and predicted prices

```
In [20]: plt.scatter(Y_train,training_data_prediction)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title("Actual prices vs Predicted Price")
plt.show()
```



```
In [21]: # prediction on Test data
test_data_prediction = lin_reg_model.predict(X_test)
```

```
In [22]: # R squared Error

error_score = metrics.r2_score(Y_test,test_data_prediction)
print("R squared Error : ",error_score)
```

R squared Error : 0.836576671502687

```
In [23]: plt.scatter(Y_test,test_data_prediction)
plt.xlabel("Actual Price")
```

```
plt.ylabel("Predicted Price")
plt.title("Actual prices vs Predicted Price")
plt.show()
```



2.Lasso Regression

In [24]: *# Loading the Lasso regression model*

```
lass_reg_model = Lasso()
```

In [25]: `lass_reg_model.fit(X_train,Y_train)`

Out[25]: Lasso()

Model Evaluation

In [26]: *# Prediction on Training data*

```
training_data_prediction = lass_reg_model.predict(X_train)
```


In [27]: *# R squared Error*

```
error_score = metrics.r2_score(Y_train,training_data_prediction)
print("R squared Error : ",error_score)
```

R squared Error : 0.8427856123435794

Visualize the actual prices and predicted prices

```
In [28]: plt.scatter(Y_train,training_data_prediction)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title("Actual prices vs Predicted Price")
plt.show()
```



```
In [29]: # prediction on Test data
test_data_prediction = lass_reg_model.predict(X_test)
```

```
In [30]: # R squared Error

error_score = metrics.r2_score(Y_test,test_data_prediction)
print("R squared Error : ",error_score)
```

R squared Error : 0.8709167941173195

```
In [ ]:
```